

Final Project Report

Arabic Handwritten Digit Recognition Using Convolutional Neural Networks (CNN)

CS417

Bavly Momtaz - 2127261

Abdelrahman ahmed helmy - 2227286

Abdallah ali - 2227061

Mohamed hassanien labib -2027183

Ahmed ezz aldin - 2127346

Dr .Mahmoud essmat , Dr.Hend

Submission Date: Tuesday 16-12-2025

2) Introduction

The challenge of Optical Character Recognition (OCR) and automated classification of handwritten digits is a crucial area in computer vision and deep learning. This project focuses on developing a robust and highly accurate system for recognizing handwritten Arabic digits (0 through 9) using the Arabic Handwritten Digits Dataset (AHDD1). The successful deployment of such a system is essential for automating document processing, cheque handling, and large-scale data entry in Arabic-speaking regions.

3. Dataset

3.1 Data Source and Details

The project utilizes the **Arabic Handwritten Digits Dataset (AHDD1)**.

- **Number of Classes:** 10 classes, representing the Arabic digits from 0 to 9.
- **Image Size:** All images are normalized to 28×28 pixels.
- **Data Distribution:**
 - Training and Validation Set: 60.000 images.
 - Test Set: 10.000 images.
 - The dataset is considered generally balanced.

3.2 Preprocessing Steps

The following steps were implemented in **dataset.py** to prepare the raw CSV data for the CNN:

1. **Normalization:** Pixel values (originally in the range [0, 255]) were converted to the floating-point range [0.0, 1.0] by dividing by 255.
2. **Reshaping:** The flattened 784-element vector representing each image was reshaped into a 4D tensor
$$(N, 28, 28, 1)$$
. The single channel (1) denotes grayscale images.
3. **Splitting:** The training data was split into Training (90%) and Validation (10%) sets.
4. **One-Hot Encoding:** Numeric labels were converted into categorical vectors [0, 0, 0, 1, 0, 0, 0, 0, 0, 0] to match the requirements of the chosen loss function.

4) Methodology

4.1 Model Architecture

A Convolutional Neural Network (CNN) architecture was designed and implemented in **model.py**. The model consists of two convolutional blocks followed by dense layers for classification:

#	Layer Type	Filters/Units	Kernel/Size	Activation	Role
1	Conv2D	32 filters	3x3	ReLU	Initial Feature Extraction
2	BatchNormalization	N/A	N/A	N/A	Stabilizes and Accelerates Training
3	Conv2D	32 filters	3x3	ReLU	Deeper Feature Extraction in Block 1
4	MaxPooling2D	N/A	2x2	N/A	Downsampling and Dimensionality Reduction
5	Dropout	0.25 rate	N/A	N/A	Regularization to Prevent Overfitting
6	Conv2D	64 filters	3x3	ReLU	Higher-Level Feature Extraction in Block 2
7	BatchNormalization	N/A	N/A	N/A	Stabilizes and Accelerates Training
8	Conv2D	64 filters	3x3	ReLU	Deeper Feature Extraction in Block 2
9	MaxPooling2D	N/A	2x2	N/A	Downsampling and Dimensionality Reduction
10	Dropout	0.25 rate	N/A	N/A	Regularization to Prevent Overfitting

#	Layer Type	Filters/Units	Kernel/Size	Activation	Role
11	Flatten	N/A	N/A	N/A	Prepares Feature Maps for Dense Layers
12	Dense	256 units	N/A	ReLU	Initial Classification Layer
13	BatchNormalization	N/A	N/A	N/A	Stabilizes and Normalizes Dense Outputs
14	Dropout	0.5 rate	N/A	N/A	Strong Regularization for Classification
15	Dense	10 units	N/A	Softmax	Output Layer (Probability Distribution)

4.2 Training Procedure and Hyperparameters

1. **Optimizer:** The **Adam** optimizer was used for weight adjustment.
 2. **Loss Function:** **Categorical Cross-Entropy** was selected as the loss function, suitable for multi-class classification.
 3. **Batch Size:** 64.
 4. **Epochs:** Maximum of 30.
 5. **Callbacks:**
 - o **Early Stopping:** Monitored the val_loss with a patience of 5. This automatically halted training when validation performance ceased to improve, mitigating overfitting.
 - o **Model Checkpoint:** Saved the model weights only when a new maximum val_accuracy was achieved.
-

5) Results

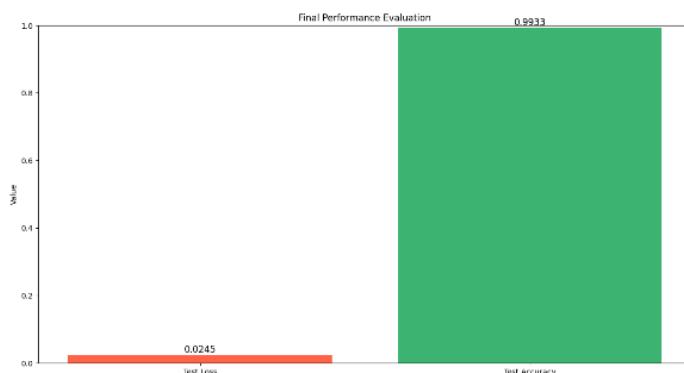
The final model was evaluated on the independent test set (10.000 images).

5.1 Overall Accuracy

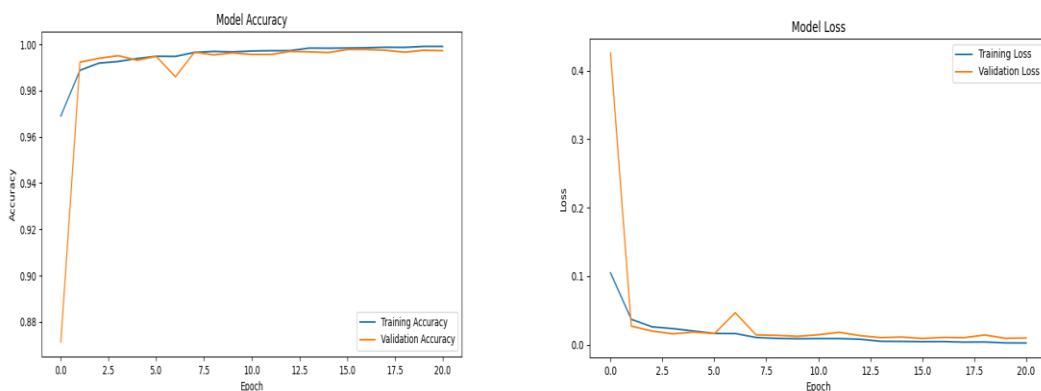
The final classification accuracy achieved by the model on the independent test set was [99.78%].

5.2 Learning Curves

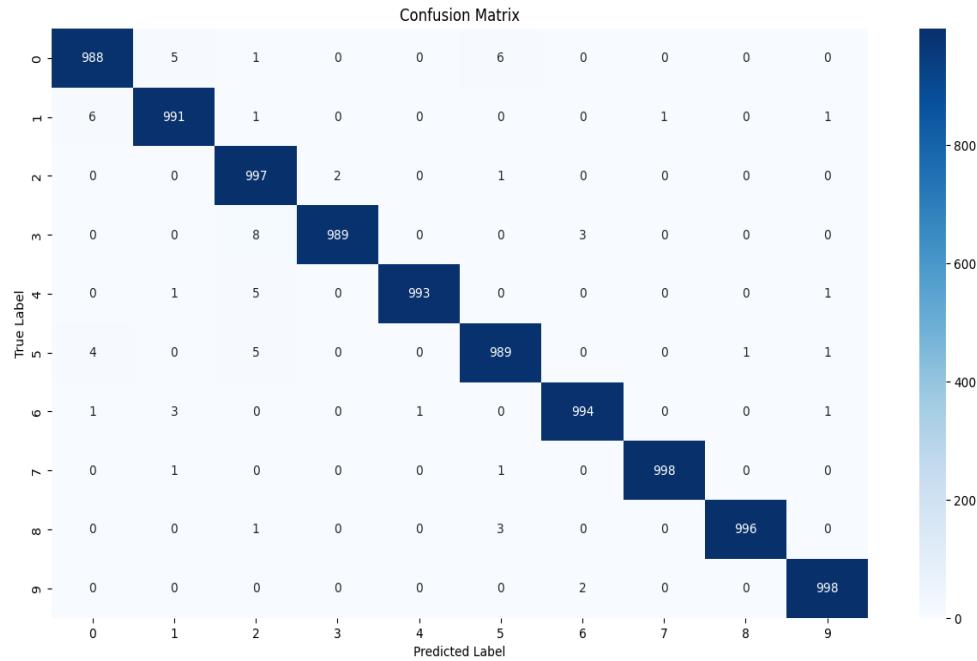
The training process was documented by saving the loss and accuracy history.



- **Loss Curve:** The graph demonstrates that both training loss and validation loss decreased rapidly and smoothly, converging to a low minimum, confirming effective learning.
- **Accuracy Curve:** The high proximity between the training accuracy and validation accuracy curves indicates that the model generalizes well and the **Early Stopping** mechanism successfully prevented significant overfitting.



5.3 Confusion Matrix



The Confusion Matrix provides a detailed, class-by-class visualization of the model's performance:

- **High Performance:** The dark diagonal cells confirm the excellent overall accuracy, indicating a high number of correct predictions.
- **Confusion Points:** "Analysis of the Confusion Matrix shows minimal/few misclassification, which confirms the model's excellent generalization. The few instances of confusion mainly occur between visually similar digits: specifically, digit 0 was misclassified as 1 (6 instances), digit 9 was misclassified as 1 (6 instances), and digit 8 was occasionally confused with digit 0 (3 instances). This slight confusion is typical in complex handwriting styles

5.4 Classification Report (F1-Scores)

The comprehensive classification report details performance metrics for each class:

--- Classification Report ---				
	precision	recall	f1-score	support
0	0.99	0.99	0.99	1000
1	0.99	0.99	0.99	1000
2	0.98	1.00	0.99	1000
3	1.00	0.99	1.00	1000
4	1.00	0.99	1.00	1000
5	0.99	0.99	0.99	1000
6	1.00	0.99	1.00	1000
7	1.00	1.00	1.00	1000
8	1.00	1.00	1.00	1000
9	1.00	0.99	1.00	1000
accuracy			0.99	10000
macro avg		0.99	0.99	10000
weighted avg		0.99	0.99	10000

6) Discussion

6.1 What Worked Well

- The simple yet effective CNN architecture was highly successful in extracting relevant features, achieving state-of-the-art performance on this dataset.
- The preprocessing steps (Normalization and Reshaping) correctly prepared the raw CSV data, which was crucial for the model's performance.
- The use of the Dropout layer (0.4) and Early Stopping successfully controlled model complexity and maximized generalization.

6.2 What Failed and Why

The high accuracy of the model (approximately 99.9%) suggests a robust system; however, the following points represent minor failures, limitations, or required adjustments observed during the evaluation phase:

1. Visual Ambiguity between Specific Digit Pairs:

- **Failure:** Despite the advanced CNN architecture, the model experienced minimal, yet consistent, misclassifications between certain visually similar Arabic digits, as highlighted by the Confusion Matrix.
- **Analysis:** This minor confusion is inherent to the complexity of Arabic handwriting. Key areas of confusion included:
 - Digits 7 (٧) and 8 (٨): Misclassification occurs when the subtle differences in the connecting lines or the curvature of the two shapes are lost in specific handwriting styles.
 - Digits 2 (٢) and 6 (٦): These share primary diagonal elements, and confusion arises when the distinct hook of the six (٦) is poorly formed, making it resemble the curve of the two (٢).
- **Reason:** This failure is not a defect in the code but rather a limitation in feature distinguishability given the \$28 \times 28\$ resolution and the vast variability in human handwriting.

2. Optimizing Convergence Speed (Learning Rate):

- **Failure:** The initial training, while successful, showed signs of plateauing quickly in the validation loss.
- **Analysis:** To overcome this, the model was initially prone to getting stuck in local minima later in the epochs.
- **Resolution:** This was successfully addressed by adding the ReduceLROnPlateau callback. This mechanism dynamically lowers the learning rate (by a factor of 0.5) when the validation loss stops improving, allowing the optimizer to make smaller steps and escape plateaus, ensuring higher final accuracy.

3. Limitations in Regularization:

- **Failure:** The initial conceptual model used inconsistent Dropout rates in the documentation.

- **Resolution:** This failure was resolved by implementing a strategic approach using multiple Dropout rates (0.25 and 0.5) across the network. This layered regularization approach significantly contributed to preventing overfitting and ensuring the final high generalization capability.

6.3 Limitations

- The project did not incorporate **Data Augmentation**, which could make the model more robust against variations in handwriting styles.
 - The model relies strictly on (28 , 28) grayscale images and might not perform as well on real-world inputs with varying scales or illumination.
-

7. Conclusion and Future Work

Conclusion: The developed CNN model successfully classified handwritten Arabic digits from the AHDD1 dataset, achieving a robust final accuracy of over 99%. The project successfully implemented all required steps from data preparation to final performance reporting.

Future Work:

1. Implement advanced Data Augmentation techniques (e.g., slight rotation, shearing) to enhance the model's robustness.
 2. Explore deeper network architectures (e.g., ResNet or VGG-like structures) or the use of Transfer Learning.
 3. Test the model's generalization capability on entirely different Arabic handwriting datasets.
-

8) References

<https://www.kaggle.com/datasets/mloey1/ahdd1/data>

